Lab 3: Out-of-Order Pipeline with In-Order Commit (10 Pts)
Part A Due: Thursday, October 14, 2021 (11:59 pm ET)
Part B (+C) Due: Sunday, October 24, 2021 (11:59 pm ET)
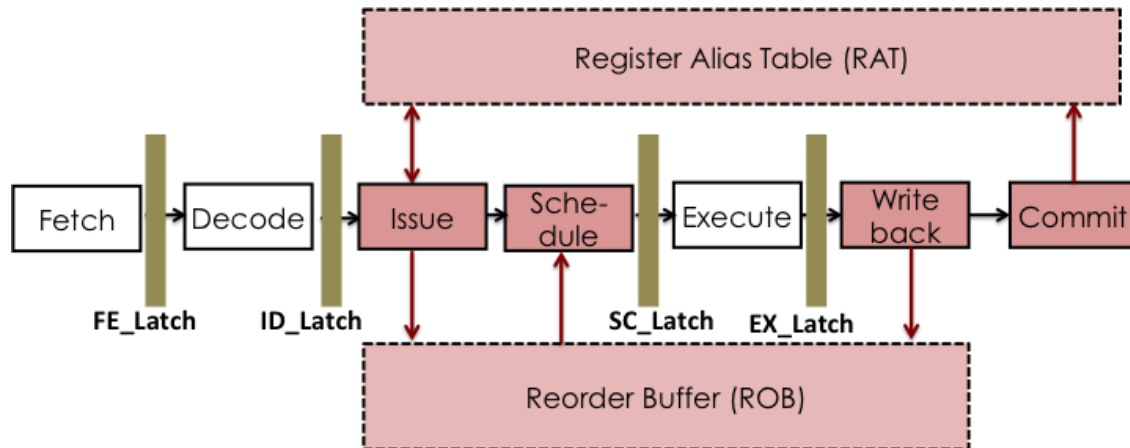
All boxes in Red have to be implemented by students for the Lab.

This is an individual assignment. You can discuss this assignment with other classmates but you should code your assignment individually. You are NOT allowed to see the code of (or show your code to) other students.

OBJECTIVE
The objective of the third programming assignment is to do a performance evaluation of an out-of-order machine. In particular, you will equip the code to do register renaming, implement different scheduling algorithms (in order and out-of-order), and use a ROB to maintain a precise state. You will initially implement a 1-wide pipeline and later extend it to a 2-wide superscalar pipeline.

PROBLEM DESCRIPTION

A seven stage out-of-order pipeline is shown in the Figure above. It consists of Fetch, Decode, Issue, Schedule, Execute, Writeback, and Commit stages. The newly added units include the Register Alias Table (RAT) and the ReOrder Buffer (ROB). *More details about these individual units and newly added stages can be found in the handout for this lab.* To simplify this assignment, we will assume perfect branch prediction. We will ignore the cc_read and cc_write operations generated by the instructions (otherwise, you would need to rename the cc register also). We will also assume that memory instructions (LD/ST) in the pipeline do not conflict with each other, to avoid the complexity of the hardware associated with memory disambiguation.

We will assume that all instructions, except LD, incur a latency of 1 cycle to execute. The latency of the LD instructions can be varied using configuration parameters.

We will use a trace driven simulator that is strictly meant for doing timing simulation. To keep the framework simple, we will not be doing any functional simulation -- which means the trace records that is fed to the pipelined machine does not contain any data values, and your pipeline will not track any data values (in REG, Memory, ROB, PC) either. Furthermore, the traces only contain the committed path instructions. The purpose of our simulation is to figure out how many clock cycles it takes to execute the given instruction stream, for a variety of different machines such as different scheduling policies, LD latencies, and pipeline width.

You will be provided with a trace reader, as well as a pipeline machine for which the fetch() decode() and exe() stage are already filled for you. Your job is to do the following:

Part A (1 point): Create the functionality for the two newly added units RAT and ROB. We have already created the *.h files for these objects, and your job is to fill the corresponding functions in the *.cpp file. The objective of this part is simply to create the three objects and compile them independently (using g++ -c filename.cpp). Note that you will not be able to debug these structures just yet, as you will need to have a working pipeline for testing. So, for this part, as long as you fill the *.cpp and submit the two files that can compile without error you will receive 1 point. You will be able to change these files for Part B. Note that if you do not do this part with seriousness, you may find Part B impossible to finish within a few days!

Part B (9 points): For a 1-wide machine, integrate the created objects in your pipeline and populate the four functions of the pipeline: issue(), schedule(), writeback(), and commit(). You will implement two scheduling policies: in-order and out-of-order (oldest ready first)

What experiments to run:
B.1 Schedule in-order, load latency=1 cycle
   ./sim -schedpolicy 0 -loadlatency 1
B.2 Schedule ooo-oldest first, load latency=1 cycle
   ./sim -schedpolicy 1 -loadlatency 1
B.3 Schedule in-order, load latency=4 cycle
   ./sim -schedpolicy 0 -loadlatency 4
B.4 Schedule ooo-oldest first, load latency=4 cycle
   ./sim -schedpolicy 1 -loadlatency 4

Part C (Extra Credit: 2 points): Extend your ooo machine to be 2-wide superscalar. Run the same four experiments as in part B, but for a 2-wide machine. (Add -pipewidth 2 to command line.)

WHAT to SUBMIT (on Canvas):
    A. For Part A, submit rat.cpp and rob.cpp (note you are allowed to change only the *.cpp files and not the *.h files).
    B. For Part B and Part C, you need to submit src.tar.gz and report.txt

REFERENCE MACHINE:
We will use **ece-linlabsrv01.ece.gatech.edu** as the reference machine for this course. (http://www.ece-help.gatech.edu/labs/unix/names.html).

Before submitting your code ensure that your code compiles on this machine, and generates the desired output (without any extra printf statements). Please follow the submission instructions. If you do not follow the submission file names, you will not receive the full credit. (You will not be penalized if your file names have automatic numeric suffixes added by Canvas, e.g., rat-1.cpp, rob-2.cpp, etc.)

FAQ:

    1. **What are the convention changes from Lab2 to Lab3**
        We will use a struct called inst_info to track both the ISA state (src1_reg, src2_reg, dest_reg) as well as microarchitectural state (src1_tag, src2_tag, dr_tag, src1_ready, src2_ready) as well as simulation metadata (inst_num). You can use inst_num to estimate the age of the instruction. Furthermore, we will use "-1" to denote invalid values or when a value is not needed. For example, if src1_reg=-1 it means src1_reg is not needed. This simplifies the number of variables the inst info needs to carry throughout the pipeline.

    2. **Why are we not simulating condition code?**
        It will significantly increase the complexity of the pipeline as you would need to rename the cc for each instruction that does a a cc_write. So, in essence an instruction would have two sources and two destinations (destreg and cc). To keep the simulation model tractable (so that students can finish this assignment in a couple of weeks), we will ignore cc read and cc write (in fact, we are ignoring branch instructions altogether, and they are treated as "OP_OTHER" instructions).